

ENGG5781 Matrix Analysis and Computations

Lecture 0: Overview

Wing-Kin (Ken) Ma

2016–2017 Term 2

Department of Electronic Engineering
The Chinese University of Hong Kong

Course Information

General Information

- Instructor: Wing-Kin Ma
 - office: SHB 323
 - e-mail: wkma@ee.cuhk.edu.hk
- Lecture hours and venue:
 - Thursday 11:30am–12:15pm, Lai Chan Pui Ngong Lecture Theatre, Y. C. Liang Hall (LPN-LT)
 - Friday 2:30pm–4:15pm, Room 404, William M. W. Mong Engineering Building (ERB 404).
- Class website: <http://www.ee.cuhk.edu.hk/~wkma/engg5781>

Course Contents

- This is a foundation course on matrix analysis and computations, which are widely used in many different fields, e.g.,
 - machine learning, computer vision,
 - systems and control, signal and image processing, communications, networks,
 - optimization, and many more...
- **Aim:** covers matrix analysis and computations at an advanced or research level.
- **Scope:**
 - basic matrix concepts, subspace, norms,
 - linear least squares
 - eigendecomposition, singular value decomposition, positive semidefinite matrices,
 - linear system of equations, LU decomposition, Cholesky decomposition
 - pseudo-inverse, QR decomposition
 - (advanced) tensor decomposition, advanced matrix calculus, compressive sensing, non-negative matrix factorization

Learning Resources

- Notes by the instructor will be provided.
- Recommended readings:
 - Gene H. Golub and Charles F. van Loan, *Matrix Computations* (Fourth Edition), John Hopkins University Press, 2013.
 - Roger A. Horn and Charles R. Johnson, *Matrix Analysis* (Second Edition), Cambridge University Press, 2012.
 - Jan R. Magnus and Heinz Neudecker, *Matrix Differential Calculus with Applications in Statistics and Econometrics* (Third Edition), John Wiley and Sons, New York, 2007.
 - Giuseppe Calafiore and Laurent El Ghaoui, *Optimization Models*, Cambridge University Press, 2014.
 - ECE 712 Course Notes by Prof. Jim Reilly, McMaster University, Canada (friendly notes for engineers)

http://www.ece.mcmaster.ca/faculty/reilly/ece712/course_notes.htm

Assessment and Others

- Assessment:
 - Assignments: 60%
 - * may contain MATLAB questions
 - * no late submissions would be accepted, except for exceptional cases.
 - Final examination: 40%
- Academic honesty: Students are strongly advised to read the University's guideline on academic honesty (<http://www.cuhk.edu.hk/policy/academichonesty/>).
- Sitting in is welcome, and please send me your e-mail address to keep you updated with the course.

A Glimpse of Topics

Least Squares (LS)

- **Problem:** given $\mathbf{A} \in \mathbb{R}^{m \times n}$, $\mathbf{y} \in \mathbb{R}^n$, solve

$$\min_{\mathbf{x} \in \mathbb{R}^n} \|\mathbf{y} - \mathbf{A}\mathbf{x}\|_2^2,$$

where $\|\cdot\|_2$ is the Euclidean norm; i.e., $\|\mathbf{x}\|_2 = \sqrt{\sum_{i=1}^n |x_i|^2}$.

- widely used in science, engineering, and mathematics
- assuming a tall and full-rank \mathbf{A} , the LS solution is uniquely given by


$$\mathbf{x}_{\text{LS}} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}.$$

Application Example: Linear Prediction (LP)

- let $y[0], y[1], \dots$ be a time series.
- **Model** (autoregressive (AR) model):


$$y[n] = a_1 y[n-1] + a_2 y[n-2] + \dots + a_L y[n-L] + w[n], \quad n = 0, 1, 2, \dots$$

for some coefficients $\{a_i\}_{i=1}^L$, where $w[n]$ is noise or modeling error.



- **Problem:** estimate $\{a_i\}_{i=1}^L$ from $\{y[n]\}$; can be formulated as LS
- **Applications:** time-series prediction, speech analysis and coding, spectral estimation...

A Toy Demo: Predicting Hang Seng Index




blue— Hang Seng Index during a certain time period.

red— training phase; the line is $\sum_{l=1}^L a_l y[n-l]$; \mathbf{a} is obtained by LS; $L = 10$.

green— prediction phase; the line is $\hat{y}[n] = \sum_{l=1}^L a_l \hat{y}[n-l]$; the same \mathbf{a} in the training phase.

A Real Example: Real-Time Prediction of Flu Activity



Tracking influenza outbreaks by ARGO — a model combining the AR model and Google search data.
Source: [\[Yang-Santillana-Kou2015\]](#).

Eigenvalue Problem

- **Problem:** given $\mathbf{A} \in \mathbb{R}^{n \times n}$, find a $\mathbf{v} \in \mathbb{R}^n$ such that

$$\mathbf{A}\mathbf{v} = \lambda\mathbf{v}, \quad \text{for some } \lambda.$$

- **Eigendecomposition:** let $\mathbf{A} \in \mathbb{R}^{n \times n}$ be symmetric; i.e., $a_{ij} = a_{ji}$ for all i, j . It also admits a decomposition


$$\mathbf{A} = \mathbf{Q}\mathbf{\Lambda}\mathbf{Q}^T,$$

where $\mathbf{Q} \in \mathbb{R}^{n \times n}$ is orthogonal, i.e., $\mathbf{Q}^T\mathbf{Q} = \mathbf{I}$; $\mathbf{\Lambda} = \text{Diag}(\lambda_1, \dots, \lambda_n)$

- also widely used, either as an analysis tool or as a computational tool
- no closed form in general; can be numerically computed


Application Example: PageRank

- PageRank is an algorithm used by Google to rank the pages of a search result.
- the idea is to use counts of links of various pages to determine pages' importance.



Source: Wiki.

One-Page Explanation of How PageRank Works



- Model:

$$\sum_{j \in \mathcal{L}_i} \frac{v_j}{c_j} = v_i, \quad i = 1, \dots, n,$$

where c_j is the number of outgoing links from page j ; \mathcal{L}_i is the set of pages with a link to page i ; v_i is the importance score of page i .


- as an example,

$$\overbrace{\begin{bmatrix} 0 & \frac{1}{2} & 1 & \frac{1}{3} \\ 0 & 0 & 0 & \frac{1}{3} \\ 0 & \frac{1}{2} & 0 & \frac{1}{3} \\ 0 & 0 & 0 & 0 \end{bmatrix}}^{\mathbf{A}} \overbrace{\begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ v_4 \end{bmatrix}}^{\mathbf{v}} = \overbrace{\begin{bmatrix} v_1 \\ v_2 \\ v_3 \\ v_4 \end{bmatrix}}^{\mathbf{v}}.$$

- finding \mathbf{v} is an eigenvalue problem—with n being of the order of millions!
- further reading: [\[Bryan-Tanya2006\]](#)

Low-Rank Matrix Approximation

- **Problem:** given $\mathbf{Y} \in \mathbb{R}^{m \times n}$ and an integer $r < \min\{m, n\}$, find an $(\mathbf{A}, \mathbf{B}) \in \mathbb{R}^{m \times r} \times \mathbb{R}^{r \times n}$ such that either $\mathbf{Y} = \mathbf{AB}$ or $\mathbf{Y} \approx \mathbf{AB}$.



- **Formulation:**

$$\min_{\mathbf{A} \in \mathbb{R}^{m \times r}, \mathbf{B} \in \mathbb{R}^{r \times n}} \|\mathbf{Y} - \mathbf{AB}\|_F^2,$$

where $\|\cdot\|_F$ is the Frobenius, or matrix Euclidean, norm.

- **Applications:** dimensionality reduction, extracting meaningful features from data, low-rank modeling, ...

Application Example: Image Compression

- let $\mathbf{Y} \in \mathbb{R}^{m \times n}$ be an image.

(a) original image, size= 102×1347



ENGG 5781 Matrix Analysis and Computations

- store the low-rank factor pair (\mathbf{A}, \mathbf{B}) , instead of \mathbf{Y} .

(b) truncated SVD, $k=5$



ENGG 5781 Matrix Analysis and Computations

(c) truncated SVD, $k=10$



ENGG 5781 Matrix Analysis and Computations

(d) truncated SVD, $k=20$




ENGG 5781 Matrix Analysis and Computations

Application Example: Principal Component Analysis (PCA)

- **Aim:** given a set of data points $\{\mathbf{y}_1, \mathbf{y}_2, \dots, \mathbf{y}_n\} \subset \mathbb{R}^n$ and an integer $k < \min\{m, n\}$, perform a low-dimensional representation

$$\mathbf{y}_i = \mathbf{Q}\mathbf{c}_i + \boldsymbol{\mu} + \mathbf{e}_i, \quad i = 1, \dots, n,$$

where $\mathbf{Q} \in \mathbb{R}^{m \times k}$ is a basis; \mathbf{c}_i 's are coefficients; $\boldsymbol{\mu}$ is a base; \mathbf{e}_i 's are errors




Toy Demo: Dimensionality Reduction of a Face Image Dataset



A face image dataset. Image size = 112×92 , number of face images = 400. Each \mathbf{x}_i is the vectorization of one face image, leading to $m = 112 \times 92 = 10304$, $n = 400$.

Toy Demo: Dimensionality Reduction of a Face Image Dataset



Mean face



1st principal left
singular vector




2nd principal left
singular vector



3rd principal left
singular vector



400th left singu-
lar vector



Energy Concentration

Singular Value Decomposition (SVD)

- **SVD:** Any $\mathbf{Y} \in \mathbb{R}^{m \times n}$ can be decomposed into

$$\mathbf{Y} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T,$$

where $\mathbf{U} \in \mathbb{R}^{m \times m}$, $\mathbf{V} \in \mathbb{R}^{n \times n}$ are orthogonal; $\mathbf{\Sigma} \in \mathbb{R}^{m \times n}$ takes a diagonal form.

- also a widely used analysis and computational tool; can be numerically computed
- SVD can be used to solve the low-rank matrix approximation problem

$$\min_{\mathbf{A} \in \mathbb{R}^{m \times r}, \mathbf{B} \in \mathbb{R}^{r \times n}} \|\mathbf{Y} - \mathbf{AB}\|_F^2.$$

Linear System of Equations

- **Problem:** given $\mathbf{A} \in \mathbb{R}^{n \times n}$, $\mathbf{y} \in \mathbb{R}^n$, solve

$$\mathbf{Ax} = \mathbf{y}.$$

- **Question 1:** How to solve it?
 - don't tell me answers like $\mathbf{x} = \text{inv}(\mathbf{A}) * \mathbf{y}$ or $\mathbf{x} = \mathbf{A} \backslash \mathbf{y}$ on MATLAB!
 - this is about matrix computations
- **Question 2:** How to solve it when n is very large?
 - it's too slow to do the generic trick $\mathbf{x} = \mathbf{A} \backslash \mathbf{y}$ when n is very large
 - getting better understanding of matrix computations will enable you to exploit problem structures to build efficient solvers
- **Question 3:** How sensitive is the solution \mathbf{x} when \mathbf{A} and \mathbf{y} contain errors?
 - key to system analysis, or building robust solutions


Why Matrix Analysis and Computations is Important?

- as said, areas such as signal processing, image processing, machine learning, optimization, computer vision, control, communications, . . . , use matrix operations extensively
- it helps you build the foundations for understanding “hot” topics such as
 - sparse recovery;
 - structured low-rank matrix approximation; matrix completion.

The Sparse Recovery Problem

Problem: given $\mathbf{y} \in \mathbb{R}^m$, $\mathbf{A} \in \mathbb{R}^{m \times n}$, $m < n$, find a **sparsest** $\mathbf{x} \in \mathbb{R}^n$ such that

$$\mathbf{y} = \mathbf{A}\mathbf{x}.$$




- by sparsest, we mean that \mathbf{x} should have as many zero elements as possible.

Application: Magnetic resonance imaging (MRI)

Problem: MRI image reconstruction.




(a)



(b)

Fig. a shows the original test image. Fig. b shows the sampling region in the frequency domain. Fourier coefficients are sampled along 22 approximately radial lines. Source: [\[Candès-Romberg-Tao2006\]](#)

Application: Magnetic resonance imaging (MRI)



(c)



(d)

Fig. c is the recovery by filling the unobserved Fourier coefficients to zero. Fig. d is the recovery by a sparse recovery solution. Source: [\[Candès-Romberg-Tao2006\]](#)

Low-Rank Matrix Completion

- **Application:** recommendation systems
 - in 2009, Netflix awarded \$1 million to a team that performed best in recommending new movies to users based on their previous preference¹.
- let \mathbf{Z} be a preference matrix, where z_{ij} records how user i likes movie j .

$$\mathbf{Z} = \begin{matrix} & \text{movies} \\ \begin{matrix} \text{users} \\ \left[\begin{array}{cccccc} 2 & 3 & 1 & ? & ? & 5 & 5 \\ 1 & ? & 4 & 2 & ? & ? & ? \\ ? & 3 & 1 & ? & 2 & 2 & 2 \\ ? & ? & ? & 3 & ? & 1 & 5 \end{array} \right] \end{matrix} \end{matrix}$$

- some entries z_{ij} are missing, since no one watches all movies.
- \mathbf{Z} is assumed to be of low rank; research shows that only a few factors affect users' preferences.
- **Aim:** guess the unknown z_{ij} 's from the known ones.

¹www.netflixprize.com

Low-Rank Matrix Completion


- The 2009 Netflix Grand Prize winners used low-rank matrix approximations [Koren-Bell-Volinsky2009].
- **Formulation** (oversimplified):

$$\min_{\mathbf{A} \in \mathbb{R}^{m \times r}, \mathbf{B} \in \mathbb{R}^{r \times n}} \sum_{(i,j) \in \Omega} |z_{ij} - [\mathbf{AB}]_{i,j}|^2$$

where Ω is an index set that indicates the known entries of \mathbf{Z} .

- cannot be solved by SVD
- in the recommendation system application, it's a large-scale problem
- alternating LS may be used

Toy Demonstration of Low-Rank Matrix Completion



Left: An incomplete image with 40% missing pixels. Right: the low-rank matrix completion result.
 $r = 120$.

Nonnegative Matrix Factorization (NMF)

- **Aim:** we want the factors to be non-negative
- **Formulation:**

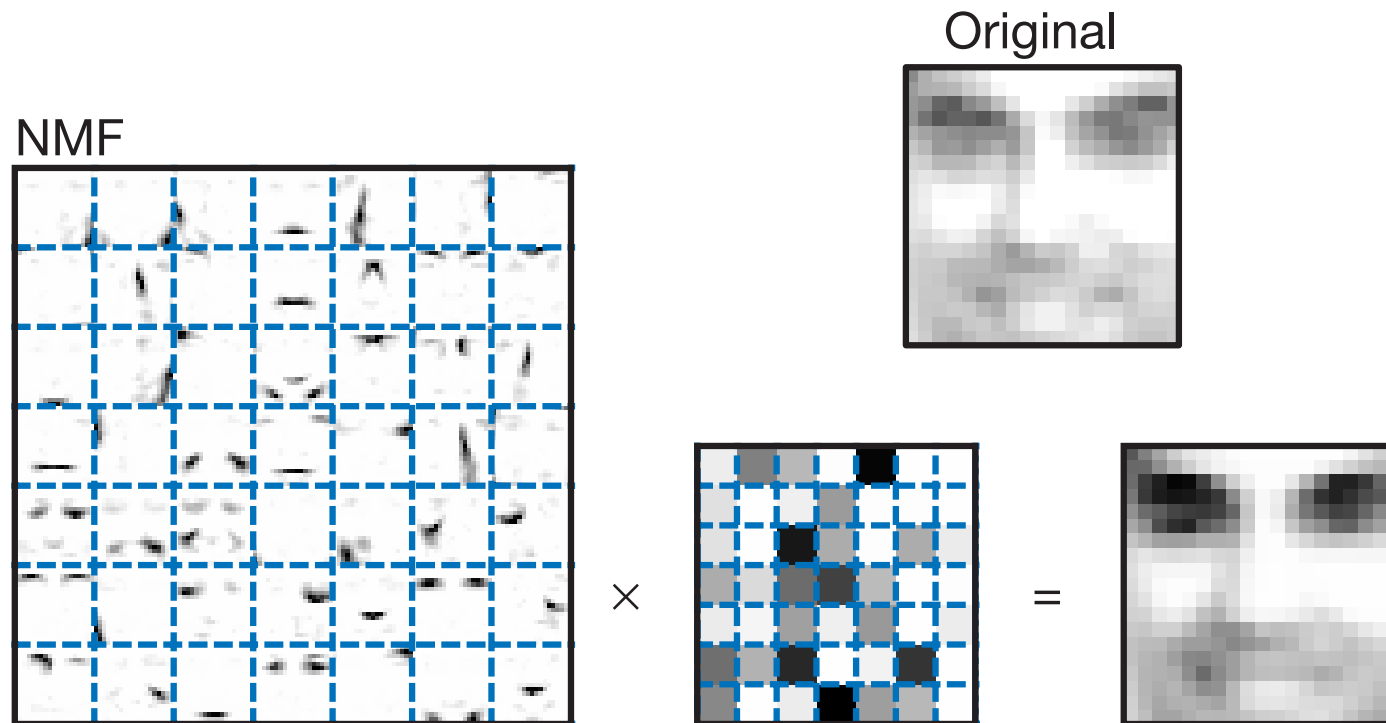
$$\min_{\mathbf{A} \in \mathbb{R}^{m \times r}, \mathbf{B} \in \mathbb{R}^{r \times n}} \|\mathbf{Y} - \mathbf{AB}\|_F^2 \quad \text{s.t. } \mathbf{A} \geq \mathbf{0}, \mathbf{B} \geq \mathbf{0},$$

where $\mathbf{X} \geq \mathbf{0}$ means that $x_{ij} \geq 0$ for all i, j .

- arguably a topic in optimization, but worth noticing
- found to be able to extract meaningful features (by empirical studies)
- numerous applications, e.g., in machine learning, signal processing, remote sensing

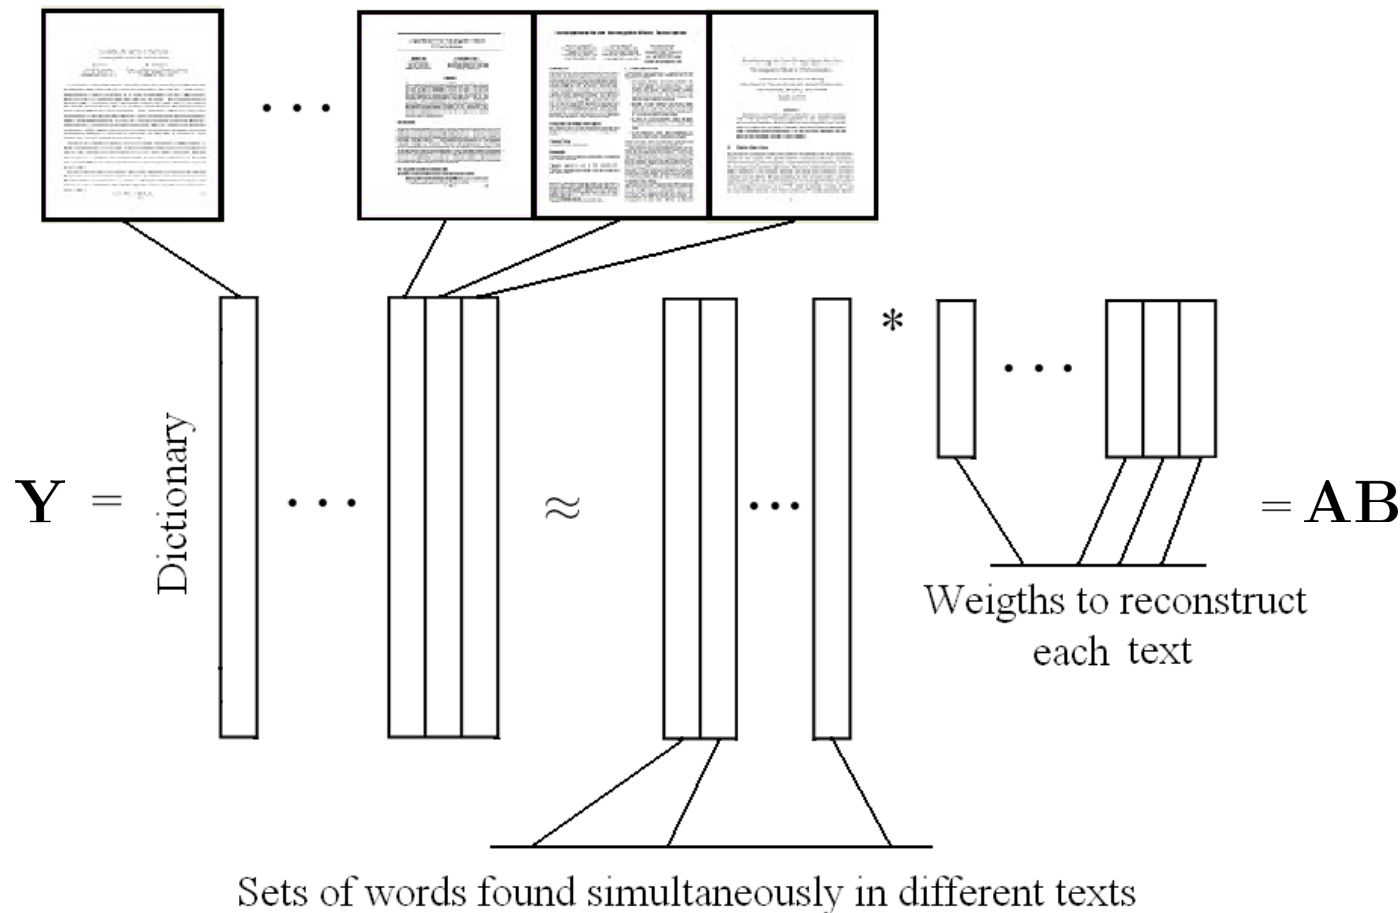
NMF Examples

- Image Processing:




The basis elements extract facial features such as eyes, nose and lips. Source: [\[Lee-Seung1999\]](#).

- **Text Mining:**




- basis elements allow us to recover different topics;
- weights allow us to assign each text to its corresponding topics.

Toy Demonstration of NMF




A face image dataset. Image size = 101×101 , number of face images = 13232. Each \mathbf{x}_i is the vectorization of one face image, leading to $m = 101 \times 101 = 10201$, $n = 13232$.

Toy Demonstration of NMF: NMF-Extracted Features



NMF settings: $r = 49$, Lee-Seung multiplicative update with 5000 iterations.


Toy Demonstration of NMF: Comparison with PCA




Mean face




1st principal left
singular vector




2nd principal left
singular vector



3th principal left
singular vector



last principal left
singular vector



Energy Concentration

A Few More Words to Say

- things I hope you will learn
 - how to read how people manipulate matrix operations, and how you can manipulate them (learn to use a tool);
 - what applications we can do, or to find new applications of our own (learn to apply a tool);
 - deep analysis skills (why is this tool valid? Can I invent new tools? Key to some topics, should go through at least once in your life time)
- feedbacks are welcome; closed-loop systems often work better than open-loop

References

- [Yang-Santillana-Kou2015]** S. Yang, M. Santillana, and S. C. Kou, “Accurate estimation of influenza epidemics using Google search data via ARGO,” *Proceedings of the National Academy of Sciences*, vol. 112, no. 47, pp. 14473–14478, 2015.
- [Bryan-Tanya2006]** K. Bryan and L. Tanya, “The 25,000,000,000 eigenvector: The linear algebra behind Google,” *SIAM Review*, vol. 48, no. 3, pp. 569–581, 2006.
- [Candès-Romberg-Tao2006]** E. J. Candès, J. Romberg, and T. Tao, “Robust uncertainty principles: Exact signal reconstruction from highly incomplete frequency information,” *IEEE Trans. Information Theory*, vol. 52, no. 2, pp. 489–509, 2006.
- [Koren-Bell-Volinsky2009]** B. Koren, R. Bell, and C. Volinsky, “Matrix factorization techniques for recommender systems,” *IEEE Computer*, vol. 42 no. 8, pp. 30–37, 2009.
- [Lee-Seung1999]** D. D. Lee and H. S. Seung, “Learning the parts of objects by non-negative matrix factorization,” *Nature*, vol. 401, no. 6755, pp. 788–791, 1999.